

[Get started](#)[Open in app](#)

Ravi Ranjan Kumar

8 Followers About

[Follow](#)

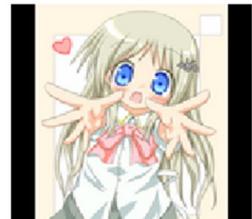
Sketch to Color Anime: An application of Conditional GAN.



Ravi Ranjan Kumar Feb 29, 2020 · 10 min read



SKETCH
to
COLOR



Don't you think it seems magical?

Introduction

On the stage of Deep Learning, things look like magical but the idea that works behind the curtain is totally conceptual. In this post we will be focusing over **Generative**

[Get started](#)[Open in app](#)

Table of contents:

1. Prerequisites
2. About Data
3. Sampling Data
4. Generative Adversarial Network (GAN)
5. Conditional Generative Adversarial Network (CGAN)
6. The Discriminator Architecture
7. The Generator Architecture
8. Different Losses
9. The GAN Model
10. Loading Data
11. Training the Model
12. Results
13. Other Applications
14. References

1. Prerequisites

This post assumes familiarity with basic knowledge of Deep Learning techniques like CNN, python syntax, some libraries like numpy, pillow, keras, etc.

2. About Data

Data has been collected from kaggle.

[Get started](#)[Open in app](#)

The Data provided is of size 10GB which contains approx 300K sketch-color pair. We will sample approx 100K sketch-color pair images for training which is approx 3GB of data. Each data provided is in form of sketch-color pair having shape (128x128x3).

3. Sampling Data

After unzipping the downloaded file, we got “danbooru-sketch-pair-128x” folder, the hierarchy tree of this folder is as follow:-

```
danbooru-sketch-pair-128x/
    color/
        sketch/
            0000/
                *.png ...
            ...
            0150/
                *.png ...
        src/
            0000/
                *.png ...
            ...
            0150/
                *.png ...
    gray/
        sketch/
            0000/
                *.png ...
            ...
            0150/
                *.png ...
        src/
            0000/
                *.png ...
            ...
            0150/
                *.png ...
```

Folder “gray” contains the gray shades images so we drop folder “gray” as we want to work only with colorful sketch-color pair.

[Get started](#)[Open in app](#)

data set, no need to keep those sub folders.

So, first we will create a dictionary having key as sub-folder name and value as list of names of images inside it, then we will sample image names from this dictionary and create our train and test set.

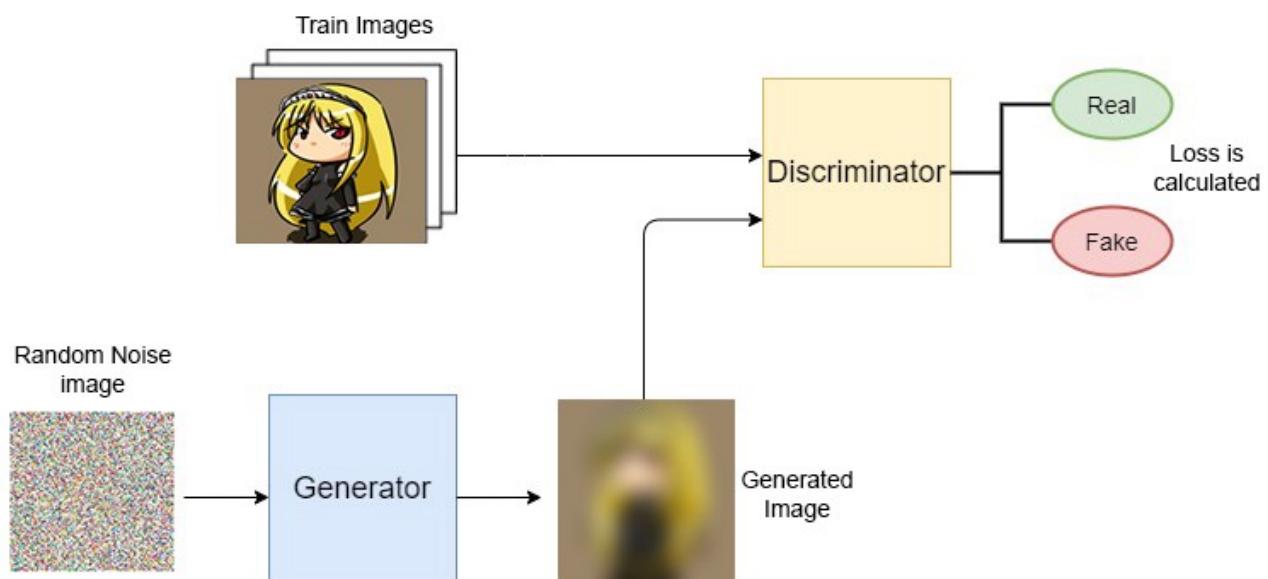
```
1  def sample_data(path):
2      """
3          This Function take a path where data is stored
4          example:- ".\danbooru-sketch-pair-128x\color\""
5          it creates two folder train_set, and test_set with sketch-color pair images
6      """
7      # storing all the names of images in a dictionary where key is subfolder name
8      # and value is a list of names of images in that sub folder
9      folder_dict = dict()
10     for d in os.listdir(path+"/sketch"):
11         folder_dict[d] = [img_name for img_name in os.listdir(path+"/sketch/"+d)]
12
13     # sampling 700 images from each sub-folder
14     sample_train = dict()
15     for key, val in folder_dict.items():
16         index = random.sample(range(0, len(val)), 700)
17         sample_train[key] = [val[i] for i in index]
18
19     # moving 700 images from each folder of sketch directory to train_set
20     split_data(sample_train, path, "sketch", "train_set")
21     # moving 700 images from each folder of src directory to train_set
22     split_data(sample_train, path, "src", "train_set")
23
24     # sampling 40 images from each folder of train_set directory to test_set
25     sample_test = dict()
26     for key, val in sample_train.items():
27         index = random.sample(range(0, 700), 40)
28         sample_test[key] = [val[i] for i in index]
29
30     # moving 40 images from each folder of sketch directory to test_set
31     split_train_data(sample_test, "train_set", "sketch", "test_set")
32     # moving 40 images from each folder of src directory to test_set
33     split_train_data(sample_test, "train_set", "src", "test_set")
34
```

[Get started](#)[Open in app](#)

Now our data is ready to process, let's see the modeling part.

4. Generative Adversarial Network (GAN)

The idea of GAN was proposed by **Ian J. Goodfellow** in june, 2014. The idea is to generate fake data after learning from the original data. GAN is a combination of two networks, the **Generator network** and the **Discriminator network**.



An Overview of GAN.

The generator takes a random noise image and tries to generate fake images similar to those which present in training set. The discriminator takes both the original image and the generated image. It tries to differentiate between real and generated image, if the output is **fake** i.e. the generated image is different from the real one then the loss is computed and the weights of discriminator is updated. This loop is repeated until we get our desired result, i.e. the discriminator should not be able to differentiate between real and generated image, in such case my generator is trained so well that it can easily fool the discriminator.

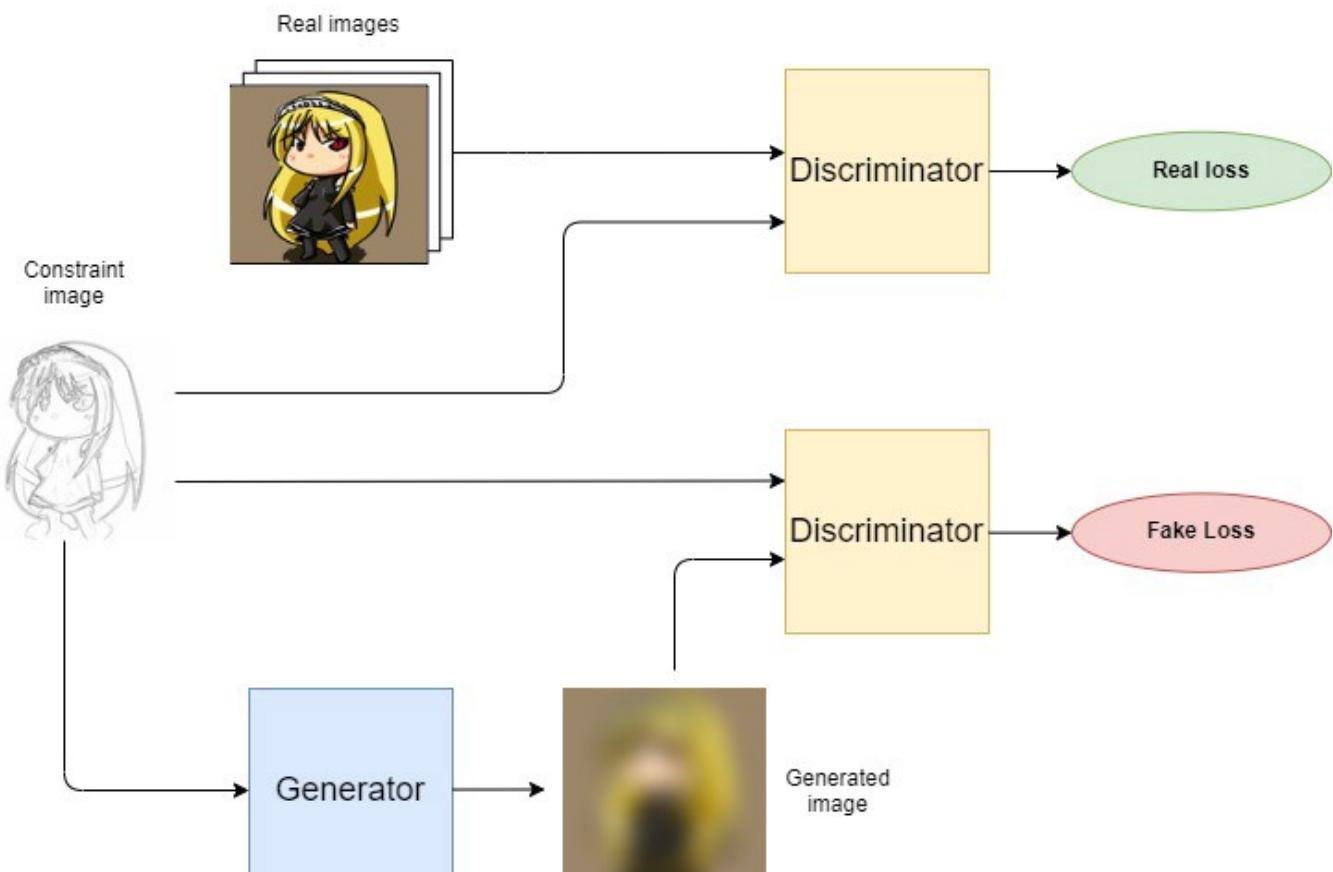
GAN comes under **Unsupervised Learning**, which uses a Supervised loss as a part of training. For understanding GAN in depth, you should read the [original GAN paper](#).

5. Conditional Generative Adversarial Network (CGAN)

[Get started](#)[Open in app](#)

constraint over generator to generate images according to the constraint provided.

CGAN also comes under **Unsupervised Learning**. I found the best reason [here](#).



An overview of Conditional GAN.

Instead of random noise vector, we are providing sketch to the generator, it is like an extra information about the distribution of training data.

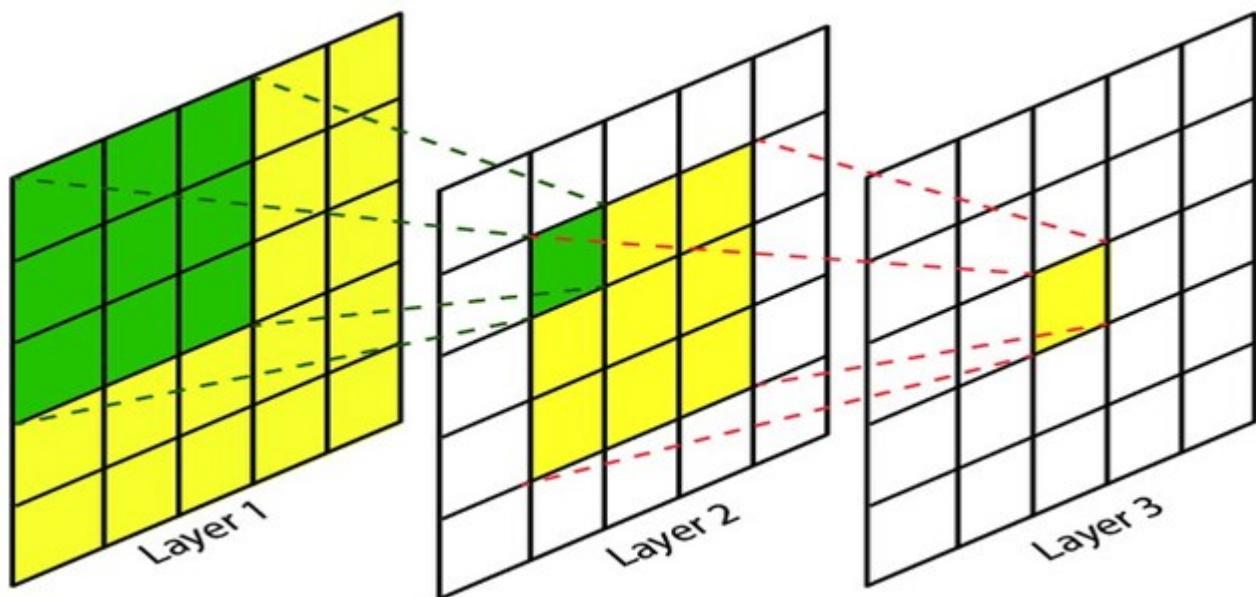
Real loss is calculated between real image and its corresponding sketch image. Fake loss is calculated between generated image and its corresponding sketch image.

Our problem falls under the category of Conditional GAN. Please go through the research paper [Image-to-Image Translation with Conditional Adversarial Networks](#) by Phillip Isola and his team, you will get a detailed idea about Conditional GAN and its wide range of applications.

Now, let's see the architecture of Discriminator and Generator in detail.

[Get started](#)[Open in app](#)

into the architecture we need to understand about **effective receptive field**. We can define it as the amount of cross-section covered by each pixel of output grid.

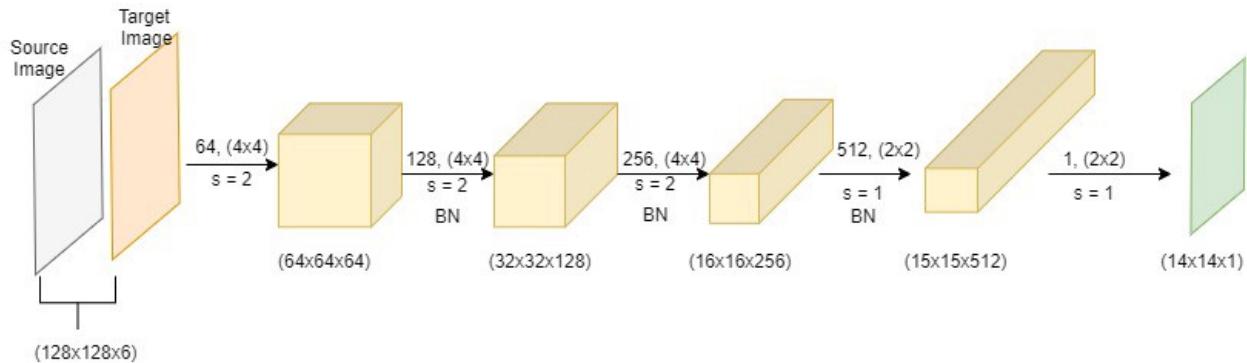


Layer 1: Input Layer | Layer 2 : Hidden Layer | Layer 3: Output Layer

Kernel (3x3) is convolved over (5x5) matrix in layer 1, therefore a single pixel in layer 2 will see (3x3) cross-section in layer 1, similarly in layer 3 a single pixel will see (3x3) in layer 2 and (5x5) in layer 1. So in layer 3 a single pixel will cover (5x5) cross section from Layer 1. This (5x5) is effective receptive field. We need to back propagate to calculate effective receptive field for a network. Formula is:-

$$\text{Receptive field} = (\text{output size} - 1) * \text{stride} + \text{kernel size}$$

Discriminator Architecture.



[Get started](#)[Open in app](#)

The **BN** is **Batch Normalization** layer, and the output is (14x14x1) grid, each cell of output layer sees (38x38) grid from the input layer. How? Each cell of output layer stores probability value that each (38x38) grid from input layer is fake or real. This technique is termed as **PatchGAN**.

```
1 def discriminator(img_shape):
2     ...
3     This function takes the shape of image as input and returns the discriminator as PatchGAN.
4     ...
5
6     # source image
7     source_image = Input(shape=img_shape)
8     # target image
9     target_image = Input(shape=img_shape)
10
11    # defining kernel_initializer
12    k_in = RandomNormal(stddev=0.02)
13
14    # concatenating images
15    con = Concatenate()([source_image, target_image])
16
17    dis = Conv2D(64, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(con) #64x64x3
18    dis = LeakyReLU(alpha=0.2)(dis)
19
20    dis = Conv2D(128, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(dis) #32x32x64
21    dis = BatchNormalization()(dis)
22    dis = LeakyReLU(alpha=0.2)(dis)
23
24
25    dis = Conv2D(256, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(dis) #16x16x128
26    dis = BatchNormalization()(dis)
27    dis = LeakyReLU(alpha=0.2)(dis)
28
29
30    # second last layer
31    dis = Conv2D(512, (2,2), kernel_initializer=k_in)(dis) #15x15x512
32    dis = BatchNormalization()(dis)
33    dis = LeakyReLU(alpha=0.2)(dis)
34
35    # last layer
36    out = Conv2D(1, (2,2), activation = "sigmoid", padding='valid', kernel_initializer=k_in)(dis)
```

[Get started](#)[Open in app](#)

```
40     # defining model
41     model = Model([source_image, target_image], final_out)
42     model.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.0002, beta_1=0.5), loss_weight
43     return model
44
45
46 # checking model
47 image_shape = (128,128,3)
48 model = discriminator(image_shape)
49
50 # summary
51 model.summary()
```

[discriminator.py](#) hosted with ❤ by GitHub[view raw](#)

7. The Generator Architecture

We are using a special architecture called U-net for Generator model. U-net is an Encoder-Decoder network with skip connections.



The U-Net Architecture | Research Paper on U-Net by Olaf Ronneberger, [click here](#) to read

The U-net takes input and tries to reduce it with a series of encoders (green color boxes in image) into a much smaller representation. The idea is that by compressing it this way we hopefully have a higher level representation of the data after the final encode layer.

[Get started](#)[Open in app](#)

U-net is used for learning image segmentation, it catches fine details about pixels and boundary regions, we also need to segment the different regions of sketch image and provide a proper color to it, this is the reason of using U-net network.

```
1 def generator(img_shape):
2
3     """
4         This function takes the shape of image as input and returns the generator model as U-net arc
5     """
6
7     k_in = RandomNormal(stddev=0.02)
8
9     # image input
10    sketch_img = Input(shape=img_shape)
11
12    # encoder
13    gen_en1 = Conv2D(64, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(sketch_i
14    gen_en1 = LeakyReLU(alpha=0.2)(gen_en1)
15
16    gen_en2 = Conv2D(128, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(gen_en1)
17    gen_en2 = BatchNormalization()(gen_en2, training=True)
18    gen_en2 = LeakyReLU(alpha=0.2)(gen_en2)
19
20    gen_en3 = Conv2D(256, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(gen_en2)
21    gen_en3 = BatchNormalization()(gen_en3, training=True)
22    gen_en3 = LeakyReLU(alpha=0.2)(gen_en3)
23
24    gen_en4 = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(gen_en3)
25    gen_en4 = BatchNormalization()(gen_en4, training=True)
26    gen_en4 = LeakyReLU(alpha=0.2)(gen_en4)
27
28    gen_en5 = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(gen_en4)
29    gen_en5 = BatchNormalization()(gen_en5, training=True)
30    gen_en5 = LeakyReLU(alpha=0.2)(gen_en5)
31
32    gen_en6 = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in)(gen_en5)
33    gen_en6 = BatchNormalization()(gen_en6, training=True)
34    gen_en6 = LeakyReLU(alpha=0.2)(gen_en6)
35
36
```

[Get started](#)[Open in app](#)

```
-->
40      # decoder
41      gen_de1 = Conv2DTranspose(512, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in
42      gen_de1 = BatchNormalization()(gen_de1, training=True)
43      gen_de1 = Dropout(0.5)(gen_de1, training=True)
44      gen_de1 = Concatenate()([gen_de1, gen_en6])
45      gen_de1 = Activation('relu')(gen_de1)
46
47      gen_de2 = Conv2DTranspose(512, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in
48      gen_de2 = BatchNormalization()(gen_de2, training=True)
49      gen_de2 = Dropout(0.5)(gen_de2, training=True)
50      gen_de2 = Concatenate()([gen_de2, gen_en5])
51      gen_de2 = Activation('relu')(gen_de2)
52
53      gen_de3 = Conv2DTranspose(512, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in
54      gen_de3 = BatchNormalization()(gen_de3, training=True)
55      gen_de3 = Dropout(0.5)(gen_de3, training=True)
56      gen_de3 = Concatenate()([gen_de3, gen_en4])
57      gen_de3 = Activation('relu')(gen_de3)
58
59      gen_de4 = Conv2DTranspose(256, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in
60      gen_de4 = BatchNormalization()(gen_de4, training=True)
61      gen_de4 = Concatenate()([gen_de4, gen_en3])
62      gen_de4 = Activation('relu')(gen_de4)
63
64      gen_de5 = Conv2DTranspose(128, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in
65      gen_de5 = BatchNormalization()(gen_de5, training=True)
66      gen_de5 = Concatenate()([gen_de5, gen_en2])
67      gen_de5 = Activation('relu')(gen_de5)
68
69      gen_de6 = Conv2DTranspose(64, (4,4), strides=(2,2), padding='same', kernel_initializer=k_in
70      gen_de6 = BatchNormalization()(gen_de6, training=True)
71      gen_de6 = Concatenate()([gen_de6, gen_en1])
72      gen_de6 = Activation('relu')(gen_de6)
73
74      # output
75      gen_col_img = Conv2DTranspose(3, (4,4), strides=(2,2), activation= 'tanh',padding='same', ke
77
78      # define model
79      model = Model(sketch_img, gen_col_img)
80      return model
```

[Get started](#)[Open in app](#)

```
84     model = generator(image_shape)
85
86     # summary
87     model.summary()
```

generator.py hosted with ❤ by GitHub

[view raw](#)

The network takes an image as input and after processing through encoder and decoder network it outputs an image.

8. Different Losses

The loss is something which we want to reduce over time with the help of optimization algorithm.

Discriminator and Generator loss

The Discriminator model tries to predict a patch from input image is real or fake. So it is trying to reduce **binary cross entropy** loss, either 0 or 1.

The loss function that is described in the research paper is:-

$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y}[\log D(x, y)] + \mathbb{E}_{x,z}[\log(1 - D(x, G(x, z)))]$$

The Conditional GAN loss. | G: Generator, D:Discriminator

‘x’ is sketch image or the constraint image, ‘y’ is real color image, ‘z’ is random image vector, In our case we are providing only ‘x’ as the input to generator to generate fake images. The Discriminator tries to **maximize** the probability of assigning the correct label to both training examples and generated samples. Generator tries to **minimize** the ‘ $\log(1 - D(x, G(x, z)))$ ’. Both models are in a tug of war situation in which both have to win by reaching an **equilibrium state**.

[Get started](#)[Open in app](#)

Generator tries to **minimize** this loss function.

To get good result over our problem we introduce more losses like L1 loss, L2 loss, Total Variation loss, let's see them in detail.

L1 loss

Popularly known as Manhattan distance. For each pixel of real color image 'y' and generated image 'G(x, z)', we find the average of distance between them . It is also known as **Pixel Level Loss**.

$$\mathcal{L}_{L1}(G) = \mathbb{E}_{x,y,z}[\|y - G(x, z)\|_1].$$

L1 loss

```
1 def pixel_loss(true, pred):
2     ...
3     parameter "true" is real color image, parameter "pred" is generated image,
4     For each pixel of real color image and generated image it returns the average
5     of distance between them.
6     ...
7     def p_loss(y_true, y_pred):
8         return K.mean( K.abs( true - pred ) )
9     return p_loss
```

[pixel_loss.py](#) hosted with ❤ by GitHub

[view raw](#)

L2 loss

Popularly known as euclidean distance. We get the feature map from the pretrained VGG 16 net and calculate the euclidean distance between feature map of real color image 'y' and generated image 'G(x, z)', It is also known as **Feature Level Loss**. We choose output of 4th layer of VGG 16 net to calculate L2 distance.

[Get started](#)[Open in app](#)

ϕ_j is feature map layer, for our case $j=4$

```
1 def feature_loss(true, pred):  
2     ...  
3     parameter "true" is feature map of real color image,  
4     parameter "pred" is feature map of generated image,  
5     it returns the euclidean distance between feature map  
6     of real color image and generated image.  
7     ...  
8     def f_loss(y_true, y_pred):  
9         return K.mean( K.sqrt( K.sum( K.square( true - pred ))))  
10    return f_loss
```

[feature_level_loss.py](#) hosted with ❤ by GitHub

[view raw](#)

Total Variation Loss

It is calculated only for generated image ' $G(x, z)$ '. It is square root of sum of differences between neighboring pixels, it is done to avoid the noise and provide smoothness in the image.

$$L_{tv} = \sqrt{(y_{i+1,j} - y_{i,j})^2 + (y_{i,j+1} - y_{i,j})^2}$$

We combine all loss together with their own weights to generate our final loss function.

```
1 def variation_loss(pred):  
2     ...  
3     parameter "pred" is generated image,  
4     it returns the square root of sum of differences between neighbouring pixels  
5     of generated image  
6     ...  
7     def v_loss(y_pred):  
8         return K.sqrt(K.sum(K.square(pred[:, 1:, :, :] - pred[:, :-1, :, :]))\n                     + K.sum(K.square(pred[:, :, 1:, :] - pred[:, :, :-1, :])))  
9     return v_loss
```

[Get started](#)[Open in app](#)

$$L = w_p L_P + w_f L_f + w_G L_G + w_{tv} L_{tv}$$

We train our model over this loss function. You can read about these losses in [this research paper](#) by Yifan Liu and his team.

9. The GAN Model

Combining both generator and discriminator results in the formation of GAN model.

While training GAN we freeze the weights of Discriminator i.e. the weights of Discriminator is not updated. This is done because we are calling both discriminator and generator in the function GAN and we can't update the weights of both network together as discriminator is using a loss function which is opposite of the generator's loss function. I got a good explanation of [this question on Quora](#).

```
1  def define_gan(g_model, d_model, image_shape):
2
3      ...
4
5      Takes generator, discriminator model and image_shape as input, and returns the GAN model as
6      ...
7
8      def gan_loss(y_true, y_pred):
9          ...
10         returns the loss function of GAN.
11
12         ...
13
14         return tf.keras.losses.binary_crossentropy(y_true, y_pred) + pixel_loss_weight * pixel_1
15             variation_loss_weight * variation_loss_out(y_pred) + feature_loss_weight * feature_
16
17
18
19     # Extracting intermediate layer features keras : https://keras.io/applications/#vgg16
20
21
22     vgg_1 = Model(inputs=vgg16.input, outputs=ReLU()(vgg16.get_layer('block2_conv2').output))
23     vgg_2 = Model(inputs=vgg16.input, outputs=ReLU()(vgg16.get_layer('block2_conv2').output))
24
25
26     pixel_loss_weight = 100
27     variation_loss_weight = .0001
28     feature_loss_weight = .01
```

[Get started](#)[Open in app](#)

```
26     sketch_image = Input(image_shape)
27     gen_output = g_model([sketch_image])
28
29     #Discriminator
30     dis_output = d_model([sketch_image, gen_output])
31
32     #Pixel Loss
33     color_image = Input(image_shape)
34     pixel_loss_out = pixel_loss(color_image, gen_output)
35
36     #Variation Loss
37     variation_loss_out = variation_loss(gen_output)
38
39     #Feature Loss
40     vgg1_out = vgg_1([tf.image.resize(color_image, (128,128), tf.image.ResizeMethod.BILINEAR)])
41     vgg2_out = vgg_2([tf.image.resize(gen_output, (128,128), tf.image.ResizeMethod.BILINEAR)])
42
43     feature_loss_out = feature_loss(vgg1_out,vgg2_out)
44
45     #Final Model
46     model = Model(inputs=[sketch_image, color_image], outputs=dis_output)
47
48     #Single output multiple loss functions in keras : https://stackoverflow.com/a/51705573/907909
49
50     model.compile(loss= gan_loss, optimizer= Adam(lr=0.0002, beta_1=0.5))
51
52     return model
53
54
55     # checking the model
56     image_shape = (128,128,3)
57
58     g_model = generator(image_shape)
59
60     d_model = discriminator(image_shape)
61
62     gan_model = define_gan(g_model, d_model,image_shape)
63     gan_model.summary()
```

[Get started](#)[Open in app](#)

“generate_real_images”, this function will read the images from folders and return the actual sketch-color pair with class label as “1” which denotes real images.

```
1 def generate_real_images(sketch_path, color_image_path, n_sample):
2     """
3         This function will read the data from paths provided and return the
4         actual sketch-color pair along with class label as "1" which denotes real images.
5     """
6
7     index = np.random.randint(0, total_img, n_sample)
8     sketch_image = []
9     color_image = []
10
11    for sketch, img in zip(sketch_path[index], color_image_path[index]):
12        sketch_image.append(np.array(Image.open(sketch).convert('RGB')))
13        color_image.append(np.array(Image.open(img).convert('RGB')))
14
15    # Normalizing the values to be between [-1, 1].
16    sketch_image = (np.array(sketch_image, dtype='float32')/127.5 - 1)
17    color_image = (np.array(color_image, dtype='float32')/127.5 - 1)
18
19    y_real = np.ones((n_sample,1))
20
21    return sketch_image, color_image, y_real
```

generate_real_images.py hosted with ❤ by GitHub

[view raw](#)

We also load fake images. The function “generate_fake_images” read only sketch image from folder and generate fake images over sketch images and return fake sketch-color pair along with class label as “0” which denotes fake images.

```
1 def generate_fake_images(g_model, sketch_path, n_sample , seed_sketch=None, seed_color=None):
2     """
3         This function will read only sketch data, and generate fake data using generator and return
4         fake sketch-color pair along with class label as "0" which denotes fake images.
5     """
6
7     sketch_image = []
8     fake_color_image = []
9
10    for sketch in sketch_path:
11        sketch_image.append(np.array(Image.open(sketch).convert('RGB')))
```

[Get started](#)[Open in app](#)

```
12         return seed_color, fake_color_image
13
14     else:
15         index = np.random.randint(0, total_img, n_sample)
16
17     for sketch in sketch_path[index]:
18         sketch_image.append(np.array(Image.open(sketch).convert('RGB')))
19
20     # Normalizing the values to be between [-1, 1].
21     sketch_image = (np.array(sketch_image, dtype='float32')/127.5 - 1)
22
23     fake_color_image = g_model.predict(sketch_image)
24     y_fake = np.zeros((n_sample, 1))
25
26     return sketch_image, fake_color_image, y_fake
```

generate_fake_images.py hosted with ❤ by GitHub

[view raw](#)

11. Training the Model

Training the model is tricky part, we randomly sample real images and fake images in batch and pass it for training over discriminator model. We need to compute real and fake loss for that we have to call discriminator twice for each generator call, this will cause my discriminator to over learn. So we call discriminator alternatively half of times and half of times together with half batch of data. This will train my discriminator model.

Now we sample real data and call GAN model over it, it will train the generator model, and here we freeze the discriminator weights. We display results and save model for each epoch.

```
1 def train(g_model, d_model, gan_model, sketch_path, color_image_path, seed_sketch, seed_color, n_
2     ...
3     This function train the model over randomly sampled batch of real images
4     and fake images. Both discriminator and generator is trained alternatively.
5     Results are displayed and generator model is saved after each epoch.
6     ...
7     batch_per_epoch = int(total_img / n_batch)
8     half_batch = int(n_batch / 2) # this is created for discriminator as we are calling it twice
```

[Get started](#)[Open in app](#)

```
12     generator_loss = []
13     discriminator_loss = []
14
15     for j in range(batch_per_epoch):
16
17         # discriminator model needs to be called twice
18         if not j%2: # train only for odd value of j, to avoid over training
19
20             # Discriminator real loss
21             real_sketch, real_col_image, y_real = generate_real_images(sketch_path, color_im-
22             d_loss_real = d_model.train_on_batch([real_sketch, real_col_image], y_real * .9)
23
24         if not j%3: # train only when j is multiple of 3. to avoid over training
25
26             # Discriminator fake loss
27             fake_sketch, fake_col_image, y_fake = generate_fake_images(g_model, sketch_path,
28             d_loss_fake = d_model.train_on_batch([fake_sketch, fake_col_image], y_fake)
29
30             total_dis_loss = d_loss_real + d_loss_fake
31
32             #GAN loss
33             real_sketch, real_col_image, y_real = generate_real_images(sketch_path, color_image_
34             gan_loss = gan_model.train_on_batch([real_sketch, real_col_image], y_real)
35
36             # saving loss per epoch
37             discriminator_loss.append(total_dis_loss)
38             generator_loss.append(gan_loss)
39
40
41             # output after 100 iter
42             if not j % 100:
43                 print("epoch>{}, {}/{}, d_real={:.3f}, d_fake={:.3f}, gan={:.3f}".format(i+1, j+
44                                         d_loss_
45
46             # writing to tensorboard
47             write_log(disc_callback, 'discriminator_loss', np.mean(discriminator_loss), i+1, (i+1)%3
48             write_log(gen_callback, 'generator_loss', np.mean(generator_loss), i+1, (i+1)%3==0)
49
50             #Summary after every epoch.
51             display.clear_output(True)
52             print('Time for epoch {} : {}'.format(i+1, datetime.now()-start2))
```

[Get started](#)[Open in app](#)

```
56
57     # saving model
58     save_model(i,g_model)
59
60     # Final summary
61     display.clear_output(True)
62     print('epoch>{}, {}/{}{}, d_real={:.3f}, d_fake={:.3f}, gan={:.3f}'.format(i+1, j+1, batch_per,
63                                         d_loss_fake, gan_lo
64     summary_save_plot(i, g_model, sketch_path, seed_sketch, seed_color, seed_color.shape[0]))
65
```

train.py hosted with ❤ by GitHub

[view raw](#)

We are training on approx 100K sketch-color pairs. Total time taken for **50 epoch** with **batch size 32** is approx **25 hours** on my system with 16 GB RAM, 6GB NVIDIA GeForce GTX 1060.

12. Results

Tensorboard Graph



[Get started](#)[Open in app](#)

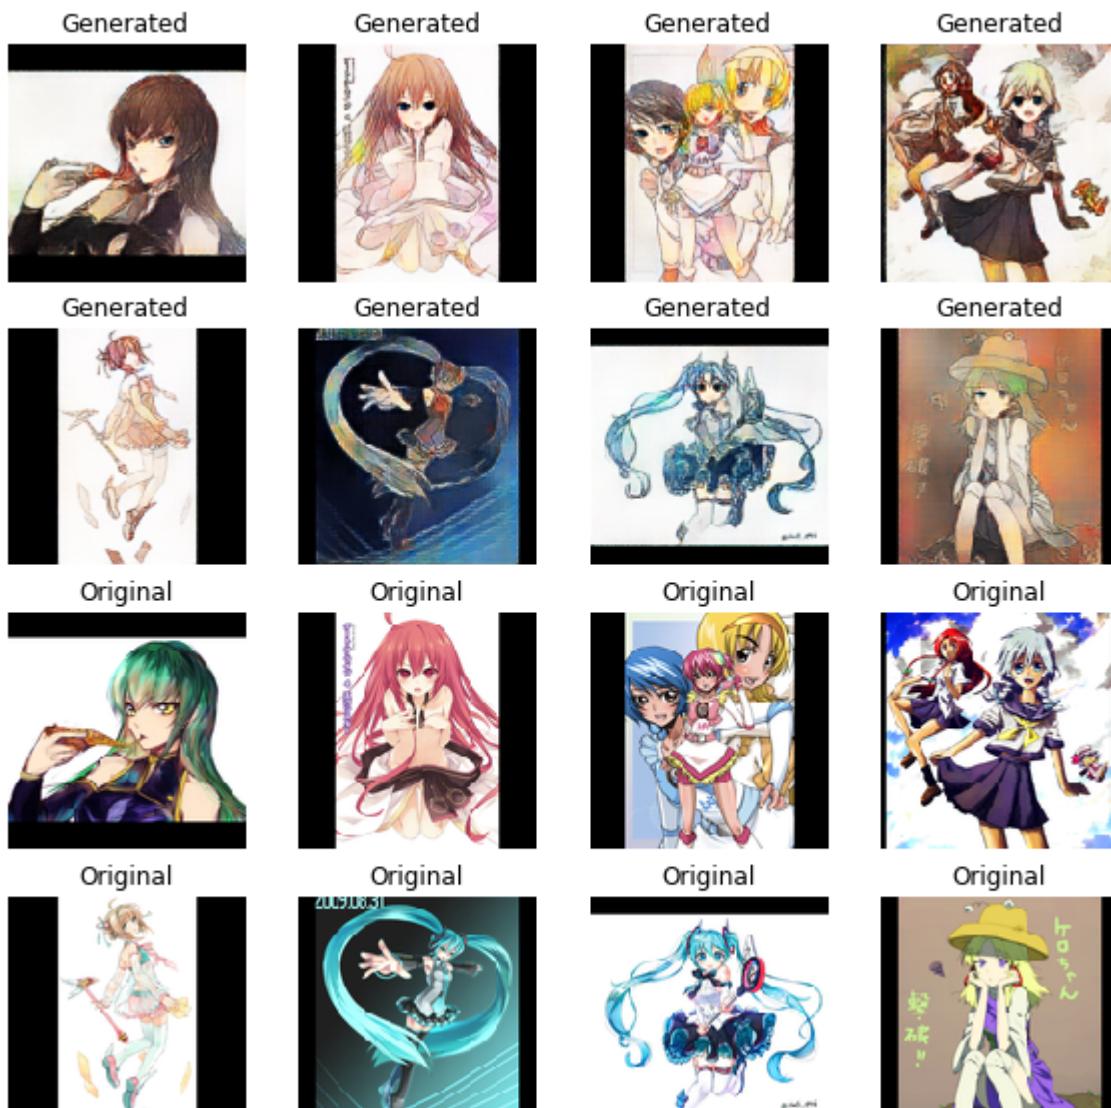
somewhat learns the distribution then the loss is increasing, and that's what we want, to increase my overall discrimination loss.

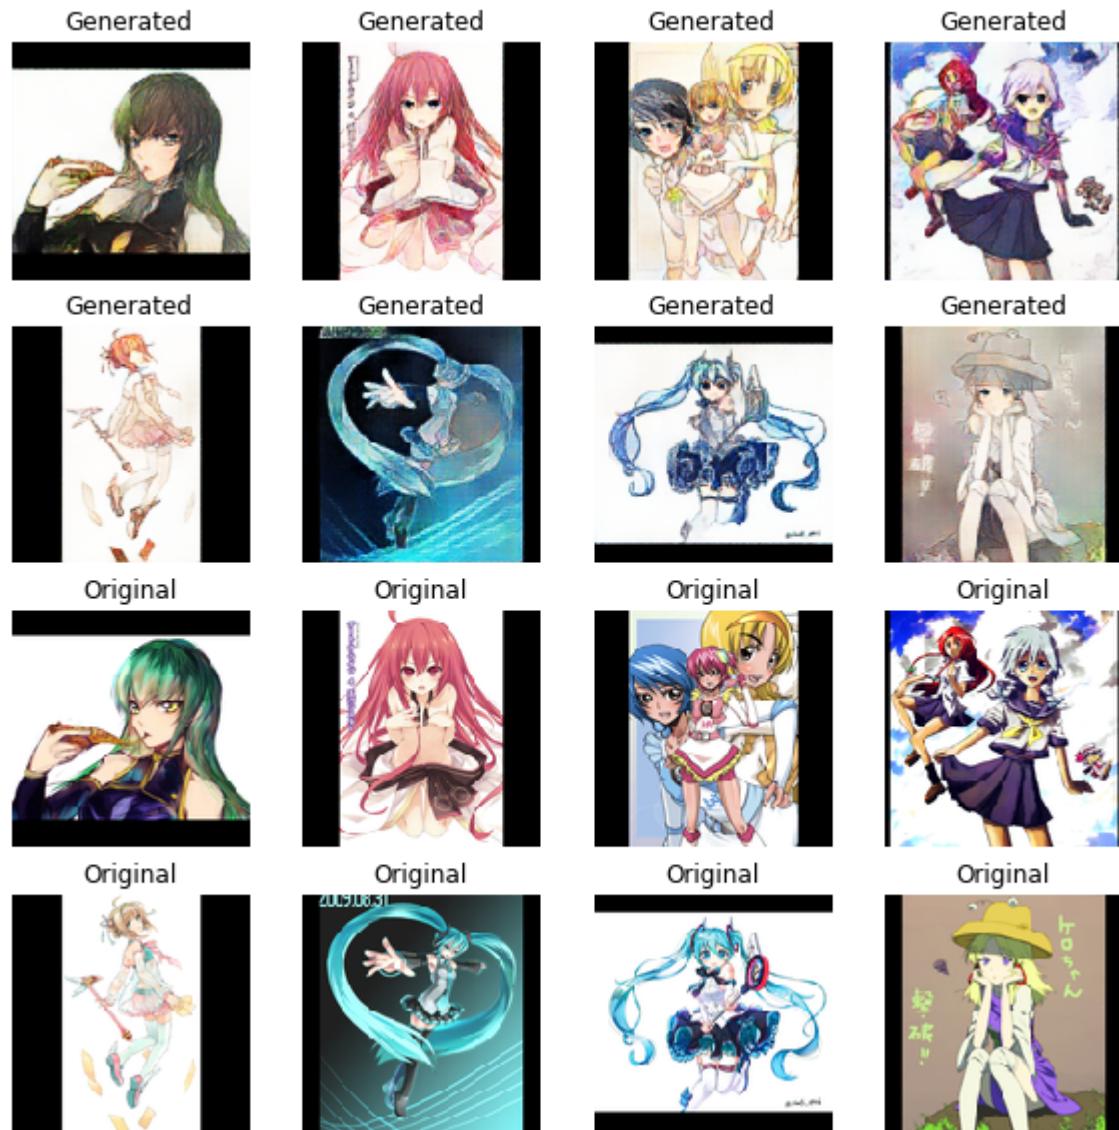
Generator loss is decreasing all way, which shows that my generator is learning all the way, which is good.

Train and test results

Now its time to check how good you are with image matching skills. Image with title '**Generated**' are fake images generated using generator and image with title '**Original**' is its corresponding real image. Let's have some work out for your eyes.

Results after different epochs from training set is:-

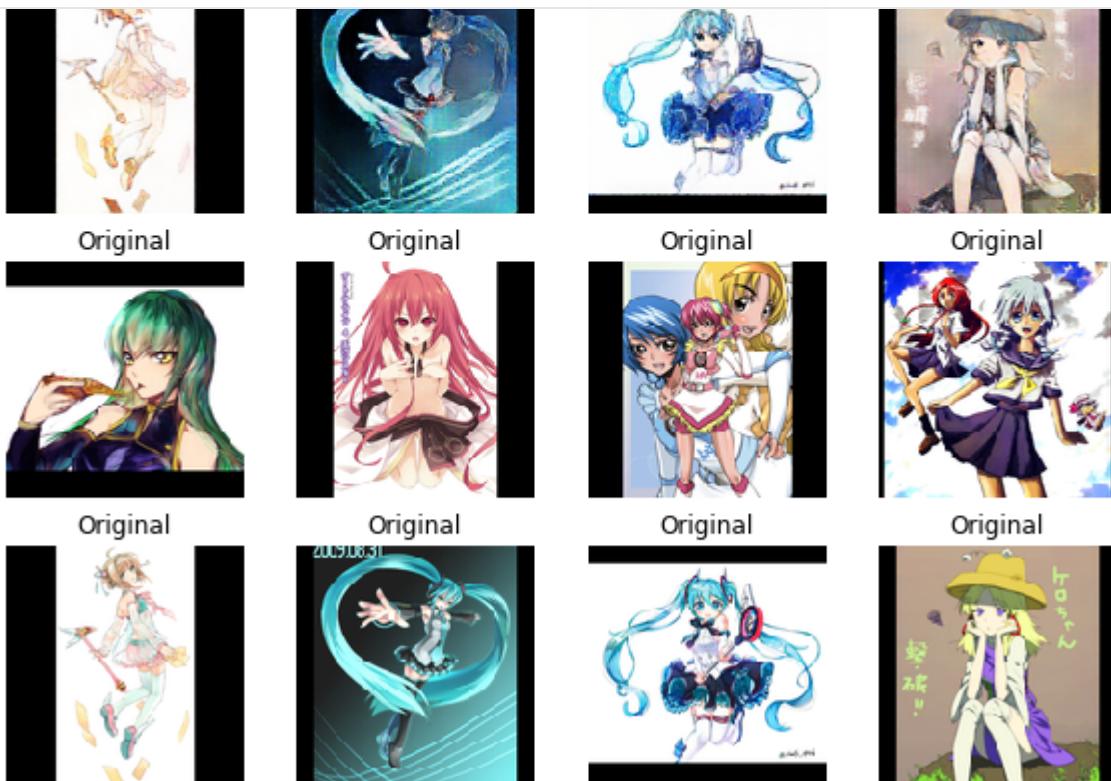


[Get started](#)[Open in app](#)

Result after 30 epochs.

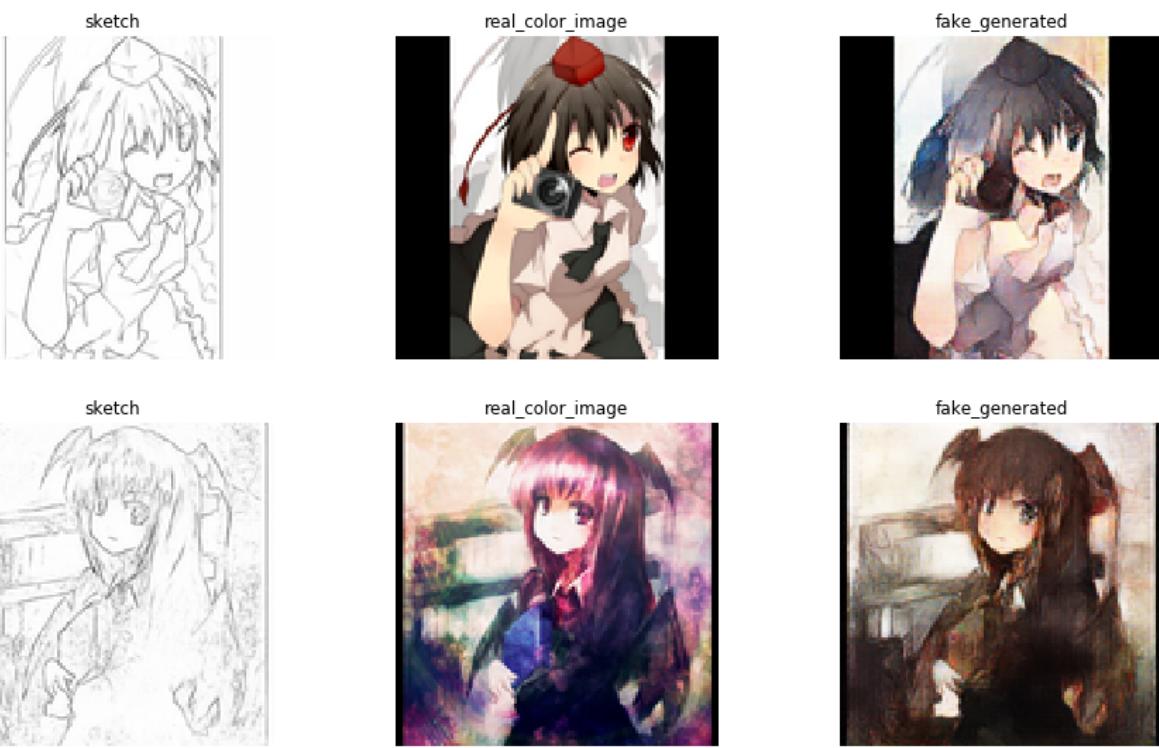
Try to observe the progress in learning as epoch increases.



[Get started](#)[Open in app](#)

Result after epoch 50

Now let's see results from the Unseen test set of data.



Get started

Open in app



sketch



fake_generated



sketch



fake_generated



sketch



fake_generated



sketch



fake_generated



Get started

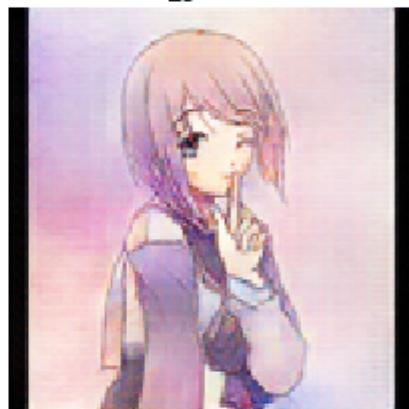
Open in app



sketch



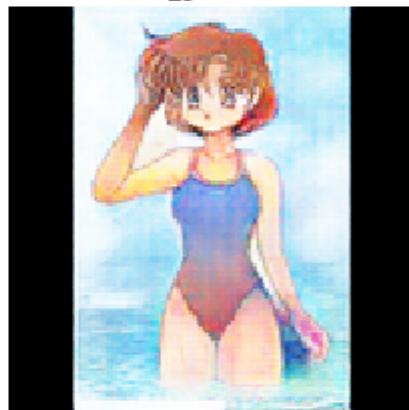
fake_generated



sketch



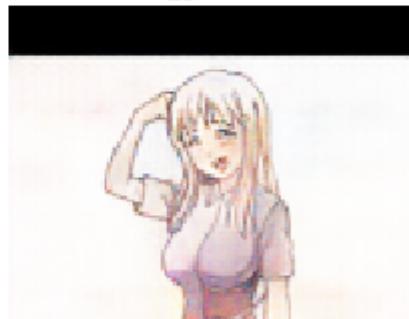
fake_generated

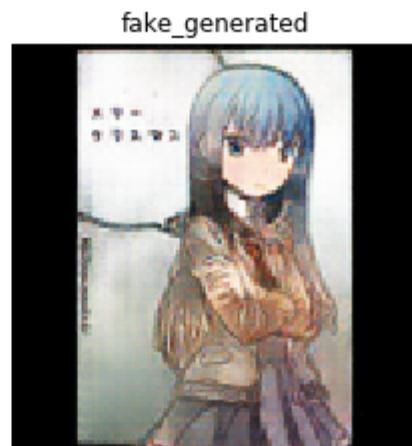
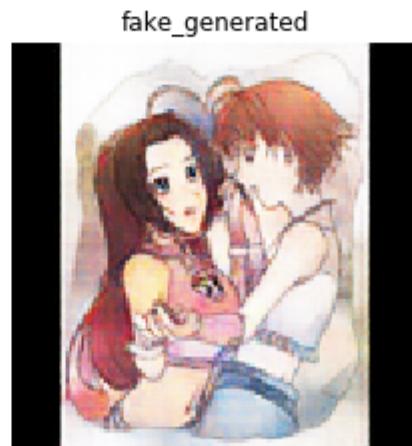
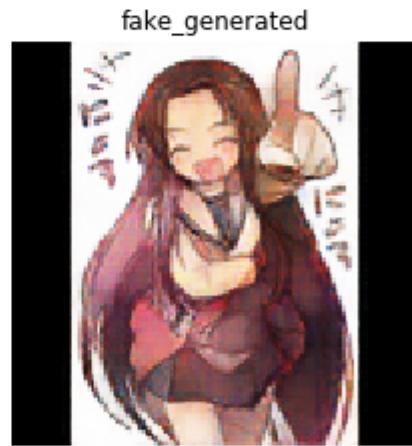


sketch

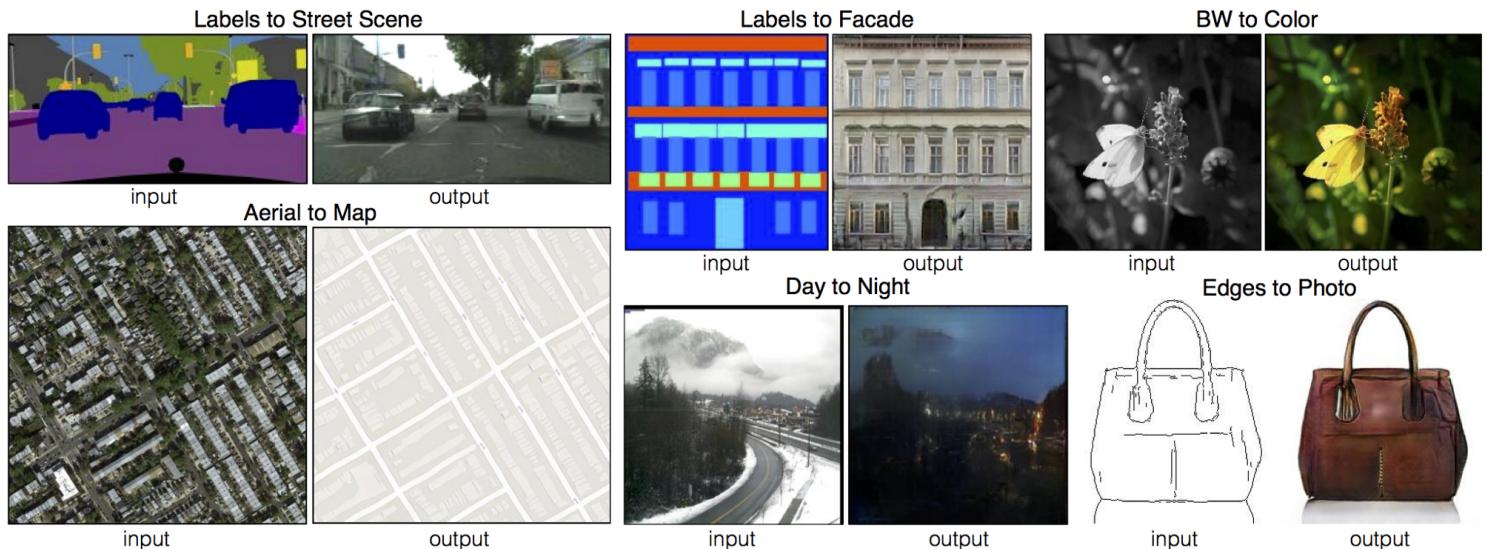


fake_generated



[Get started](#)[Open in app](#)

Result over randomly chosen image from the test set.

[Get started](#)[Open in app](#)

Source : [Image-to-Image Translation with Conditional Adversarial Nets](#)

Here input is the constraint image, and output is the desired result which is similar to ground truth image i.e. real image.

Get full [**code on github**](#) with detailed explanation in form of **jupyter notebook**.

Suggestions are welcome. Thanks a lot for reading... ❤

14. References

1. [Appliedaicourse](#)
2. Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros : [Image-to-Image Translation with Conditional Adversarial Networks](#)
3. Yifan Liu, Zengchang Qin, Zhenbo Luo, Hua Wang : [Auto-painter: Cartoon Image Generation from Sketch by Using Conditional Generative Adversarial Networks](#)
4. [Are GAN's unsupervised or supervised?](#)
5. [TensorBoard Profile: Profiling basic training metrics in Keras](#)
6. <https://phillipi.github.io/pix2pix/>

You can reach me at:-

[Get started](#)

[Open in app](#)



[edium . https://edium.com/1av1-1034005/](https://edium.com/1av1-1034005/)

Gans

Pix2pix

Generative Adversarial

Cnn

Deep Learning

About Write Help Legal

Get the Medium app

